

Better Growth Modeling: cactus case study

Tom Miller

6/21/2020

About the data

This document walks through growth modeling with the cactus (*Cylindriopuntia imbricata*, or CYIM) case study. The plants are a little oddly shaped. We measure plant size by taking maximum height, maximum width of the crown, and then an additional width measurement perpendicular to the maximum. We then convert this into the volume of a cone (cm^3) and take the natural log of this as the size variable.



There are a total of 827 individuals in the data set spanning 8 plots or spatial blocks and 9 transition-years. There are some additional, earlier years of data but they come from different plots and that was causing problems in the shrinkage model. So I filtered out these early years. Note that even in these later years, plots 7 and 8 did not start until 2011.

```
##  
##      2009 2010 2011 2012 2013 2014 2015 2016 2017  
##  1   81   72   60   67   65   64   62   79   72  
##  2   58   55   42   42   40   38   38   64   56  
##  3   81   75   70   83   84   85   80  102   99  
##  4   89   82   66   64   64   64   63  115   96  
##  5   43   43   44   47   46   45   45   49   47  
##  6   76   67   56   64   66   65   64   83   83  
##  7     0     0   47   54   53   53   51   84   79  
##  8     0     0   72   75   72   70   83   82   77
```

Gaussian fits

Use lme4 to fit Gaussian models of the form $size_{t+1} \sim size_t$. There are temporal and spatial random effects. No obvious fixed effects besides initial size (this is an observational study) but I can test whether a quadratic term for initial size improves fit. Random effects are intercept-only because convergence problems otherwise.

```
CYIM_lmer_models <- list()
CYIM_lmer_models[[1]] <- lmer(log(vol_t1) ~ log(vol_t) + (1 | year_t) + (1 | plot),
  data = CYIM, REML = F, control = lmerControl(optimizer = "bobyqa"))
CYIM_lmer_models[[2]] <- lmer(log(vol_t1) ~ log(vol_t) + I(log(vol_t)^2) + (1 | year_t) +
  (1 | plot), data = CYIM, REML = F, control = lmerControl(optimizer = "bobyqa"))
```

Now use the iterative re-weighting approach to re-fit these with non-constant variance:

```
## NegLogLik function to fit variance model for residuals
varPars = function(pars) {
  return(-sum(dnorm(resids, mean=0, sd=exp(pars[1] + pars[2]*fitted_vals), log=TRUE)))
}

pars<-list()
for(mod in 1:length(CYIM_lmer_models)) {
  err = 1; rep=0;
  ## Steve's error level was too high a bar to clear. I raised this.
  while(err > 0.0001) {
    rep=rep+1; model = CYIM_lmer_models[[mod]];
    fitted_vals = fitted(model); resids = residuals(model);
    out=optim(c(sd(resids),0),varPars,control=list(maxit=5000));
    pars[[mod]]=out$par;
    new_sigma = exp(pars[[mod]][1] + pars[[mod]][2]*fitted_vals); new_weights = 1/((new_sigma)^2)
    new_weights = 0.5*(weights(model) + new_weights); # cautious update
    new_model <- update(model,weights=new_weights);
    err = weights(model)-weights(new_model); err=sqrt(mean(err^2));
    cat(mod,rep,err,"\n") # check on convergence of estimated weights
    CYIM_lmer_models[[mod]]<-new_model;
  }
}
```

For fair AIC comparison, re-fit both models with best weights.

```
##### For a fair AIC comparison, fit all models with the same weights
aics = unlist(lapply(CYIM_lmer_models, AIC))
best_model = which(aics == min(aics))
best_weights = weights(CYIM_lmer_models[[best_model]])
for (mod in 1:length(CYIM_lmer_models)) {
  CYIM_lmer_models[[mod]] <- update(CYIM_lmer_models[[mod]], weights = best_weights)
}
aictab(CYIM_lmer_models)

## Warning in aictab.AIClmerModLmerTest(CYIM_lmer_models):
## Model names have been supplied automatically in the table

## Model selection based on AICc:
```

```

##          K      AICc Delta_AICc AICcWt Cum.Wt       LL
## Mod1 5 9016.95        0.00   0.62   0.62 -4503.47
## Mod2 6 9017.94        0.99   0.38   1.00 -4502.96

```

The linear model is favored. Last step is to re-fit with REML=T.

```

aics = unlist(lapply(CYIM_lmer_models, AIC))
best_model = which(aics == min(aics))
CYIM_lmer_best = CYIM_lmer_models[[best_model]]
best_weights = weights(CYIM_lmer_best)
CYIM_lmer_best <- lmer(log(vol_t1) ~ log(vol_t) + (1 | year_t) + (1 | plot), data = CYIM,
                        weights = best_weights, REML = TRUE)

```

Now visualize this kernel. This model is the *very best* we could do in a Gaussian framework (assuming there are no major fixed effects we are missing).

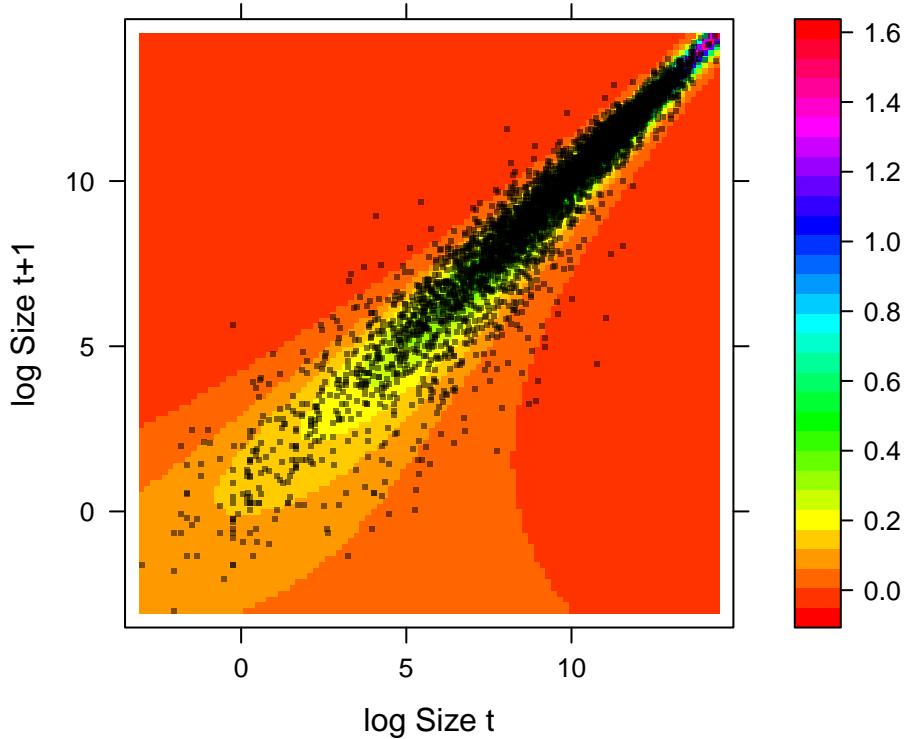
```

## dummy variable for initial size
size_dim <- 100
size_dum <- seq(min(log(CYIM$vol_t)), max(log(CYIM$vol_t)), length.out = size_dim)
## make a polygon for the Gaussian kernel with non-constant variance
CYIM_lmer_best_kernel <- matrix(NA, size_dim, size_dim)
for (i in 1:size_dim) {
  mu_size <- fixef(CYIM_lmer_best)[1] + fixef(CYIM_lmer_best)[2] * size_dum[i]
  CYIM_lmer_best_kernel[, i] <- dnorm(size_dum, mean = mu_size, sd = exp(pars[[best_model]][1] +
    pars[[best_model]][2] * mu_size))
}

levelplot(CYIM_lmer_best_kernel, row.values = size_dum, column.values = size_dum,
           cuts = 30, col.regions = rainbow(30), xlab = "log Size t", ylab = "log Size t+1",
           main = "Gaussian, non-constant variance", panel = function(...) {
             panel.levelplot(...)
             grid.points(log(CYIM$vol_t), log(CYIM$vol_t1), pch = ".", gp = gpar(cex = 3,
               col = alpha("black", 0.5)))
           })

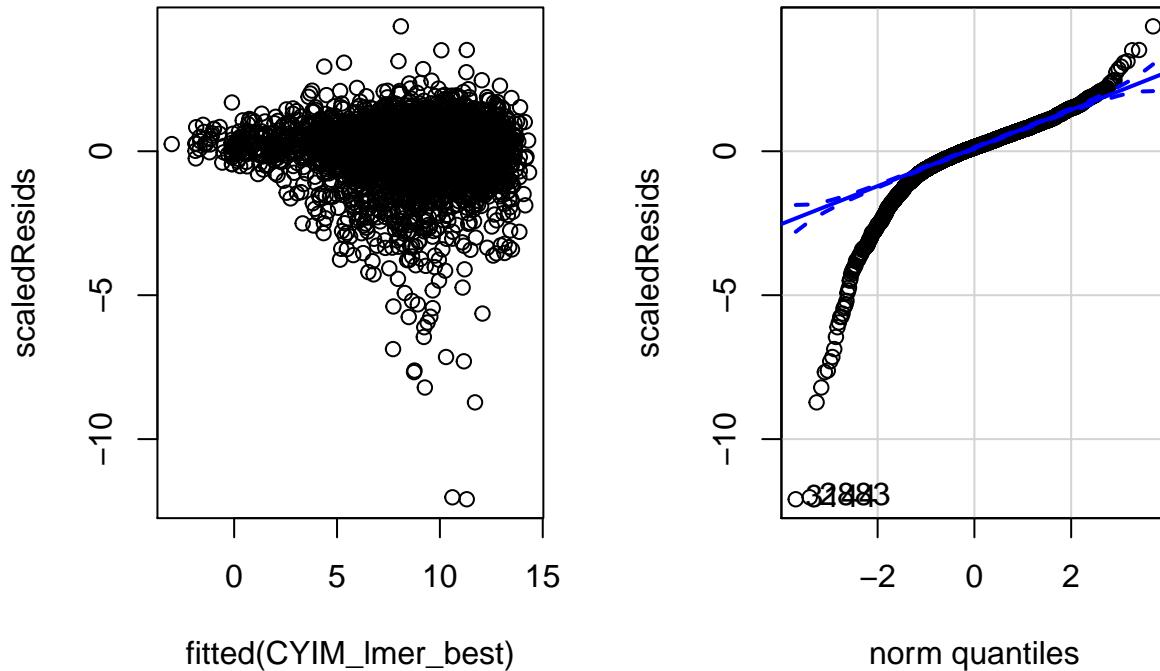
```

Gaussian, non-constant variance



Visually, this looks great. But is it any good? No. Scaled residuals are not remotely Gaussian.

```
scaledResids = residuals(CYIM_lmer_best) * sqrt(weights(CYIM_lmer_best))
par(mfrow = c(1, 2))
plot(fitted(CYIM_lmer_best), scaledResids)
qqPlot(scaledResids) # really bad in both tails
```



```
## [1] 3144 2883
```

```
jarque.test(scaledResids) # normality test: FAILS, P < 0.001
```

```
##
##  Jarque-Bera Normality Test
##
## data: scaledResids
## JB = 68502, p-value < 2.2e-16
## alternative hypothesis: greater
```

```
anscombe.test(scaledResids) # kurtosis: FAILS, P < 0.001
```

```
##
##  Anscombe-Glynn kurtosis test
##
## data: scaledResids
## kurt = 21.353, z = 32.724, p-value < 2.2e-16
## alternative hypothesis: kurtosis is not equal to 3
```

```
agostino.test(scaledResids) # skewness: FAILS, P<0.001
```

```
##
```

```

## D'Agostino skewness test
##
## data: scaledResids
## skew = -2.6615, z = -42.5617, p-value < 2.2e-16
## alternative hypothesis: data have a skewness

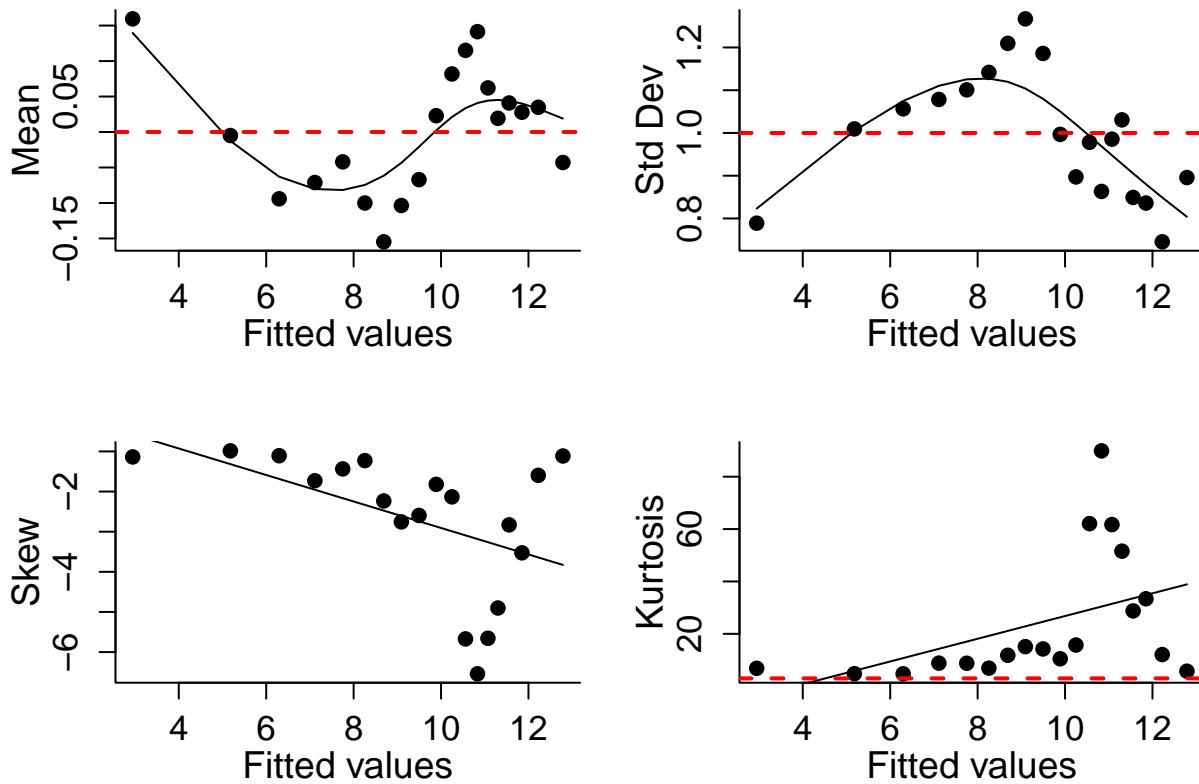
```

One last look at the standardized residuals. They are roughly mean zero and unit variance – so that checks out. But there is negative skew and excess kurtosis, especially at large sizes.

```

px = fitted(CYIM_lmer_best)
py = scaledResids
par(mfrow = c(2, 2), bty = "l", mar = c(4, 4, 2, 1), mgp = c(2.2, 1, 0), cex.axis = 1.4,
     cex.lab = 1.4)
z = rollMoments(px, py, windows = 10, smooth = TRUE, scaled = TRUE)

```



Finding a better distribution

We now know the Gaussian provides a poor fit to the residual variance. We would like to know which distribution provides a better - ideally *good* - fit. Because there is size-dependence in skew and kurtosis, we cannot marginalize over the entire distribution of residuals (this may point us in the wrong direction). Instead, we can slice up the data into bins of expected value and find the best distribution for each bin using `gamlss' fitDist()`.

```

n_bins <- 8
## I need to rewrite this code because it runs fitDist 4 separate times
select_dist <- tibble(fit_best = fitted(CYIM_lmer_best), scale_resid = residuals(CYIM_lmer_best) *
  sqrt(weights(CYIM_lmer_best))) %>% mutate(bin = as.integer(cut_number(fit_best,
  n_bins)), best_dist = NA, secondbest_dist = NA, aic_margin = NA)
for (b in 1:n_bins) {
  bin_fit <- fitDist(select_dist$scale_resid[select_dist$bin == b], type = "realline")
  select_dist$best_dist[select_dist$bin == b] <- names(bin_fit$fits[1])
  select_dist$secondbest_dist[select_dist$bin == b] <- names(bin_fit$fits[2])
  select_dist$aic_margin[select_dist$bin == b] <- bin_fit$fits[2] - bin_fit$fits[1]
}
(select_dist <- select_dist %>% group_by(bin) %>% summarise(n_bin = n(), best_dist = unique(best_dist),
  secondbest_dist = unique(secondbest_dist), aic_margin = unique(aic_margin)))

## # A tibble: 8 x 5
##   bin n_bin best_dist secondbest_dist aic_margin
##   <int> <int> <chr>      <chr>          <dbl>
## 1     1     563 SHASH    SEP4        4.01e+ 0
## 2     2     563 SHASH    SEP4        1.76e+ 0
## 3     3     562 JSU       JSUo      1.44e-10
## 4     4     563 SHASH    ST1        1.38e+ 0
## 5     5     563 ST4       ST1        3.05e- 1
## 6     6     562 ST4       NET        2.63e+ 0
## 7     7     563 ST5       ST2        1.58e- 1
## 8     8     563 ST1       ST5        1.42e- 1

```

It's a little bit of everything and convergence is not great (I've suppressed those warnings in this output). The skewed t pops up a lot though, so I will proceed with that one and hope in the end the fit is decent. To guide the skewed t fits we can visualize how the higher moments are related to the fitted value.

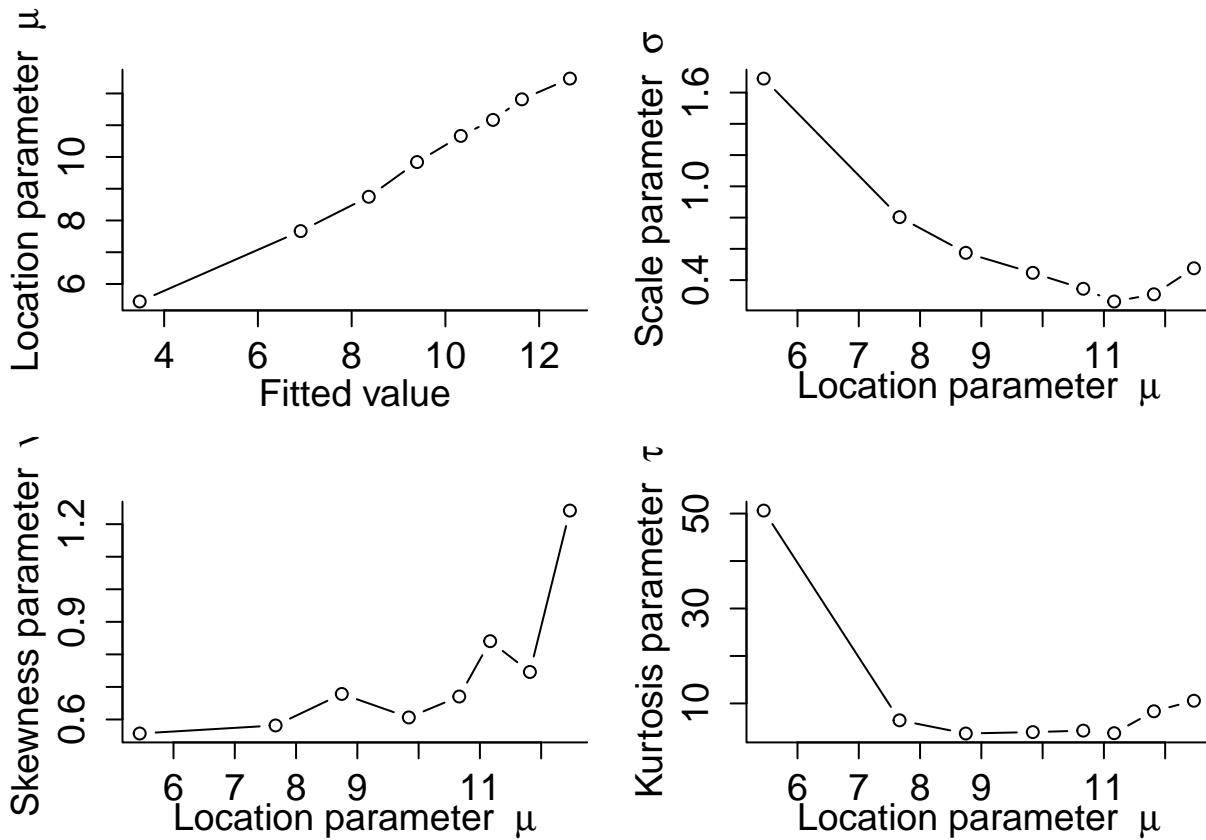
```

CYIM_bin_fit <- CYIM %>% mutate(fitted = fitted(CYIM_lmer_best), bin = as.integer(cut_number(fitted,
  n_bins))) %>% mutate(mu = NA, sigma = NA, nu = NA, tau = NA)
for (b in 1:n_bins) {
  ## I get the most stable tau estimates with ST3
  bin_fit <- gamlssML(log(CYIM_bin_fit$vol_t1[CYIM_bin_fit$bin == b]) ~ 1, family = "ST3")
  CYIM_bin_fit$mu[CYIM_bin_fit$bin == b] <- bin_fit$mu
  CYIM_bin_fit$sigma[CYIM_bin_fit$bin == b] <- bin_fit$sigma
  CYIM_bin_fit$nu[CYIM_bin_fit$bin == b] <- bin_fit$nu
  CYIM_bin_fit$tau[CYIM_bin_fit$bin == b] <- bin_fit$tau
}
CYIM_bin_fit <- CYIM_bin_fit %>% group_by(bin) %>% summarise(N = n(), mean_fitted = mean(fitted),
  mu = unique(mu), sigma = unique(sigma), nu = unique(nu), tau = unique(tau))

par(mfrow = c(2, 2), bty = "l", mar = c(4, 4, 2, 1), mgp = c(2.2, 1, 0), cex.axis = 1.4,
  cex.lab = 1.4)
## Steve's spline.scatter.smooth() function not working for me and I did not both
## trying to figure out why
plot(CYIM_bin_fit$mean_fitted, CYIM_bin_fit$mu, xlab = "Fitted value", ylab = expression(paste("Location
  mu)), type = "b")
plot(CYIM_bin_fit$mu, CYIM_bin_fit$sigma, xlab = expression(paste("Location parameter  ",
  mu)), ylab = expression(paste("Scale parameter  ", sigma)), type = "b")
plot(CYIM_bin_fit$mu, CYIM_bin_fit$nu, xlab = expression(paste("Location parameter  ",
  mu)), ylab = expression(paste("Skewness parameter  ", nu)), type = "b")

```

```
plot(CYIM_bin_fit$mu, CYIM_bin_fit$tau, xlab = expression(paste("Location parameter ", mu)), ylab = expression(paste("Kurtosis parameter ", tau)), type = "b")
```



Fitting the final model

Now we can fit a custom model via maximum likelihood, matching the structure of the best lmre model but fitting the random effects as fixed instead and estimating the corresponding variances with Steve's shrinkage methods. The likelihood function is based on the skewed t (using gamlss' dST3() - the parameterization that seemed best behaved.)

First we will define the linear predictor for the location parameter μ (which is not necessarily the expected value). Note that this with parameterization, the intercept is year1/plot1.

```
U = model.matrix(~0 + year_t + plot + log(vol_t), data = CYIM)
```

Next define a likelihood function using this linear predictor for the location. I am including quadratic terms for sigma and tau based on preliminary fit where linear models for these parameters led to a not-so-great fit at some sizes. Quadratic terms are also suggested in the moment plots above.

```
LogLik = function(pars, response, U) {
  pars1 = pars[1:ncol(U)]
  pars2 = pars[-(1:ncol(U))]
```

```

mu = U %*% pars1
val = dST3(x = response, mu = mu, sigma = exp(pars2[1] + pars2[2] * mu + pars2[3] *
    mu^2), nu = exp(pars2[4] + pars2[5] * mu), tau = exp(pars2[6] + pars2[7] *
    mu + pars2[8] * mu^2), log = T)
return(val)
}

```

Now fit it, using the lmer() fit for the mean (location) and lm() fits for the higher moments as starting parameter values. This comes straight from Steve's PSSP code, including his convergence paranoia, which is probably warranted. In preliminary fits, all the iterations converged on the same likelihoods, and this takes forever, so I am skimping on the paranoia here, but we might return to it to make this a good example to follow. This could also be a model selection step; easy to pull AIC from the maxLik fit.

```

paranoid_iter <- 1
coefs = list(paranoid_iter); LL=numeric(paranoid_iter);

# Starting values from the pilot model are jittered to do multi-start optimization.
# Using good starting values really speeds up convergence in the ML fits
# Linear predictor coefficients extracted from the lmer model
fixed_start = c(unlist(ranef(CYIM_lmer_best)$year_t) + unlist(ranef(CYIM_lmer_best)$plot)[1], #year est
               unlist(ranef(CYIM_lmer_best)$plot)[-1],
               fixef(CYIM_lmer_best)[2]) # size slope
## make sure the dimensions line up
#length(fixed_start);ncol(U);colnames(U)

# Shape and scale coefficients from the rollapply diagnostic plots
fit_sigma = lm(log(sigma)^mu + I(mu^2), data=CYIM_bin_fit)
fit_nu = lm(log(nu)^mu, data=CYIM_bin_fit)
fit_tau = lm(log(tau)^mu + I(mu^2), data=CYIM_bin_fit)
p0=c(fixed_start, coef(fit_sigma), coef(fit_nu),coef(fit_tau))

for(j in 1:paranoid_iter) {
  out=maxLik(logLik=LogLik,start=p0*exp(0.2*rnorm(length(p0))), response=log(CYIM$vol_t1),U=U,
             method="BHHH",control=list(iterlim=5000,printLevel=2),finalHessian=FALSE);

  out=maxLik(logLik=LogLik,start=out$estimate,response=log(CYIM$vol_t1),U=U,
             method="NM",control=list(iterlim=5000,printLevel=1),finalHessian=FALSE);

  out=maxLik(logLik=LogLik,start=out$estimate,response=log(CYIM$vol_t1),U=U,
             method="BHHH",control=list(iterlim=5000,printLevel=2),finalHessian=FALSE);

  coefs[[j]] = out$estimate; LL[j] = out$maximum;
  cat(j, "#-----#", out$maximum, "\n");
}

j = min(which(LL==max(LL))) ## they actually all land on the same likelihood-that's good!
out=maxLik(logLik=LogLik,start=coefs[[j]],response=log(CYIM$vol_t1),U=U,
           method="BHHH",control=list(iterlim=5000,printLevel=2),finalHessian=TRUE)

##### save results of ML fit.
names(out$estimate)<-c(colnames(U),"sigma_b0","sigma_b1","sigma_b2","nu_b0","nu_b1","tau_b0","tau_b1",
coefs=out$estimate
## these are the indices of plot and year effects, to make the next steps a little more intuitive

```

```

years=1:9
plots=10:16
SEs = sqrt(diag(vcov(out)))
AIC_ST <- 2*length(coefs) - 2*out$maximum
AIC_norm <- AIC(CYIM_lmer_best)

```

The AIC comparison is a blowout, though I am not sure if it is ok to use the lmer AIC. Might be better to re-fit with plot and year as fixed.

```
tibble(growth_function = c("Gaussian", "Skewed t"), AIC = c(AIC_ST, AIC_norm))
```

```

## # A tibble: 2 x 2
##   growth_function     AIC
##   <chr>              <dbl>
## 1 Gaussian           6986.
## 2 Skewed t           9030.

```

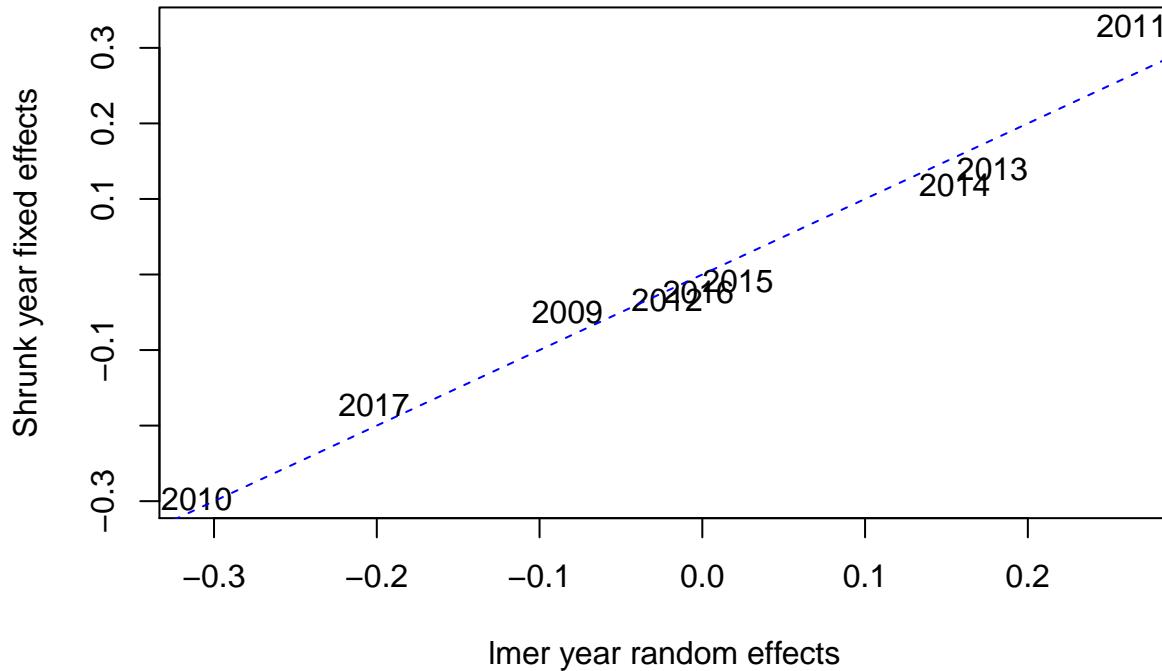
Now use Steve's shrinkage code to get the random effect variances and compare these to the lmer estimates. I will need Steve to explain to me what the shrinkage step is doing, and what is the theory for this (or I should do my own homework).

```

# shrinkage random effects for (1/year)
year_fixed.fx = coefs[years] - mean(coefs[years])
year_fixed.se = SEs[years]
year_sigma2.hat = mean(year_fixed.fx^2) - mean(year_fixed.se^2)
year_shrunk.fx = year_fixed.fx * sqrt(year_sigma2.hat/(year_sigma2.hat + year_fixed.se^2))
# lmer random effects for (1/year)
year_ran.fx = ranef(CYIM_lmer_best)[["year_t"]]

plot(year_ran.fx$year_t$(Intercept), year_shrunk.fx, xlab = "lmer year random effects",
      ylab = "Shrunk year fixed effects", type = "n")
text(year_ran.fx$year_t$(Intercept), year_shrunk.fx, labels = rownames(year_ran.fx$year_t))
abline(0, 1, col = "blue", lty = 2)

```



```
tibble(sd_estimate = c(sd(year_fixed.fx), sd(year_shrunk.fx), sd(year_ran.fx$year_t`^(Intercept)`)),
      method = c("fixed", "shrunk", "lme4"))
```

```
## # A tibble: 3 x 2
##   sd_estimate method
##       <dbl> <chr>
## 1     0.193 fixed
## 2     0.182 shrunk
## 3     0.184 lme4
```

The code for plots is more complicated than years, because plot effects were parameterized as contrasts. What is the SE of the plot 1 effect?? Is it the mean of the year SEs? – probably not! I am going to drop plot 1 from the shrunk effects, because I don't know how to calculate its SE. This will bias the variance, hopefully not too much.

These estimates don't look as good as the year effects but, again, the shrinkage variance corresponds well to the lme4 estimate.

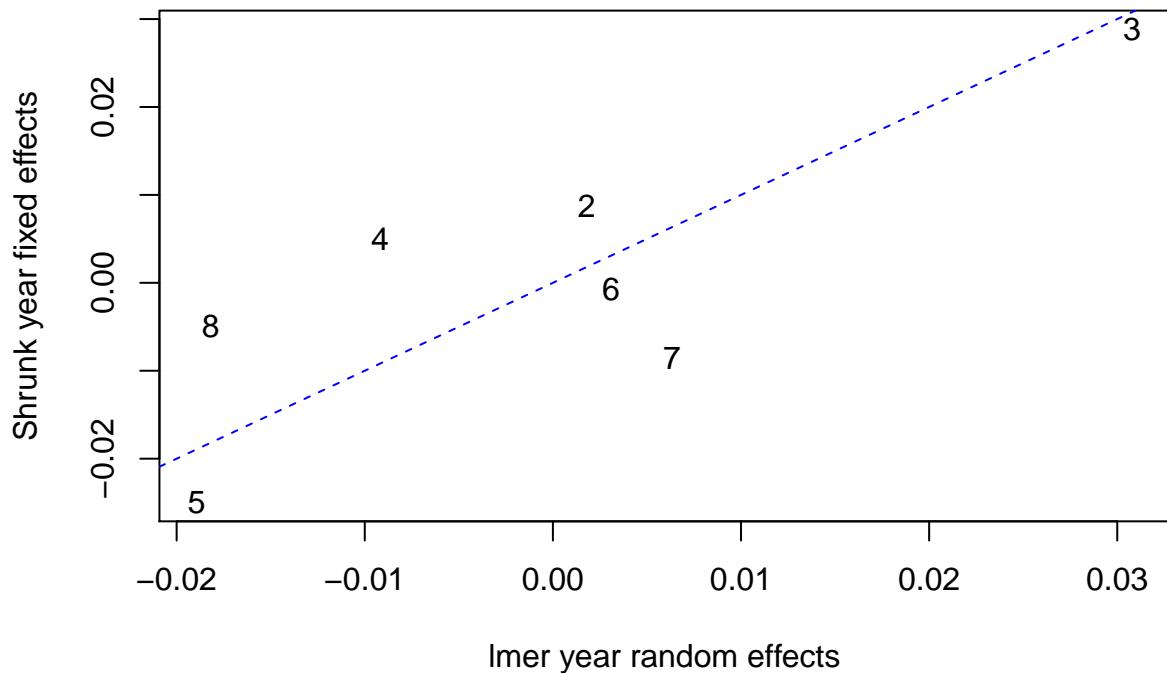
```
# shrinkage random effects for (1/plot)
plot_coefs <- mean(coefs[years]) + coefs[plots]
plot_fixed.fx = plot_coefs - mean(plot_coefs)
plot_fixed.se = SEs[plots]
plot_sigma2.hat = mean(plot_fixed.fx^2) - mean(plot_fixed.se^2)
plot_shrunk.fx = plot_fixed.fx * sqrt(plot_sigma2.hat/(plot_sigma2.hat + plot_fixed.se^2))
# lmer random effects for (1/plot)
```

```

plot_ran.fx = ranef(CYIM_lmer_best)["plot"]

plot(plot_ran.fx$plot$`(`Intercept)`), c(NA, plot_shrunk.fx), xlab = "lmer year random effects",
      ylab = "Shrunk year fixed effects", type = "n")
text(plot_ran.fx$plot$`(`Intercept)`), c(NA, plot_shrunk.fx), labels = rownames(plot_ran.fx$plot))
abline(0, 1, col = "blue", lty = 2)

```



```

tibble(sd_estimate = c(sd(plot_fixed.fx), sd(plot_shrunk.fx), sd(plot_ran.fx$plot$`(`Intercept)`))),
      method = c("fixed", "shrunk", "lme4"))

## # A tibble: 3 x 2
##   sd_estimate method
##       <dbl> <chr>
## 1     0.0303 fixed
## 2     0.0166 shrunk
## 3     0.0160 lme4

```

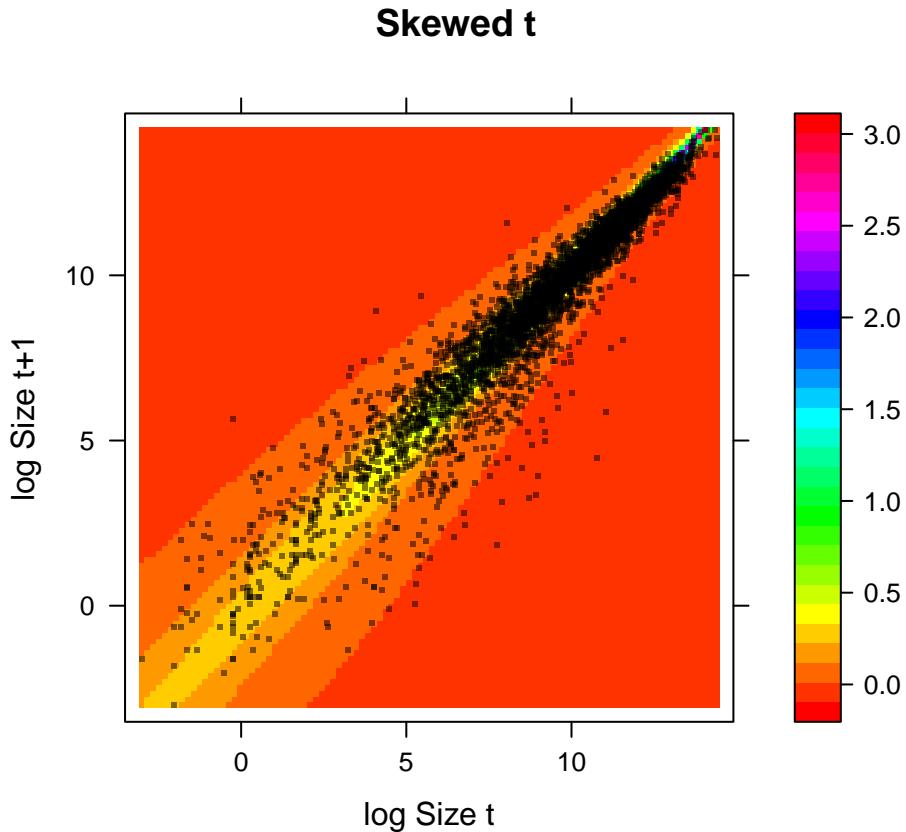
Steve had an interesting idea for an alternative way to fit two random effects, with one used as an offset for the other. Might swing back to that, especially since the reference-level problem is nagging at me.

Visual diagnostics of model fit via shrinkage

Here is the top-level view of the skewed t kernel. This will be the mean kernel, averaged over years and plots.

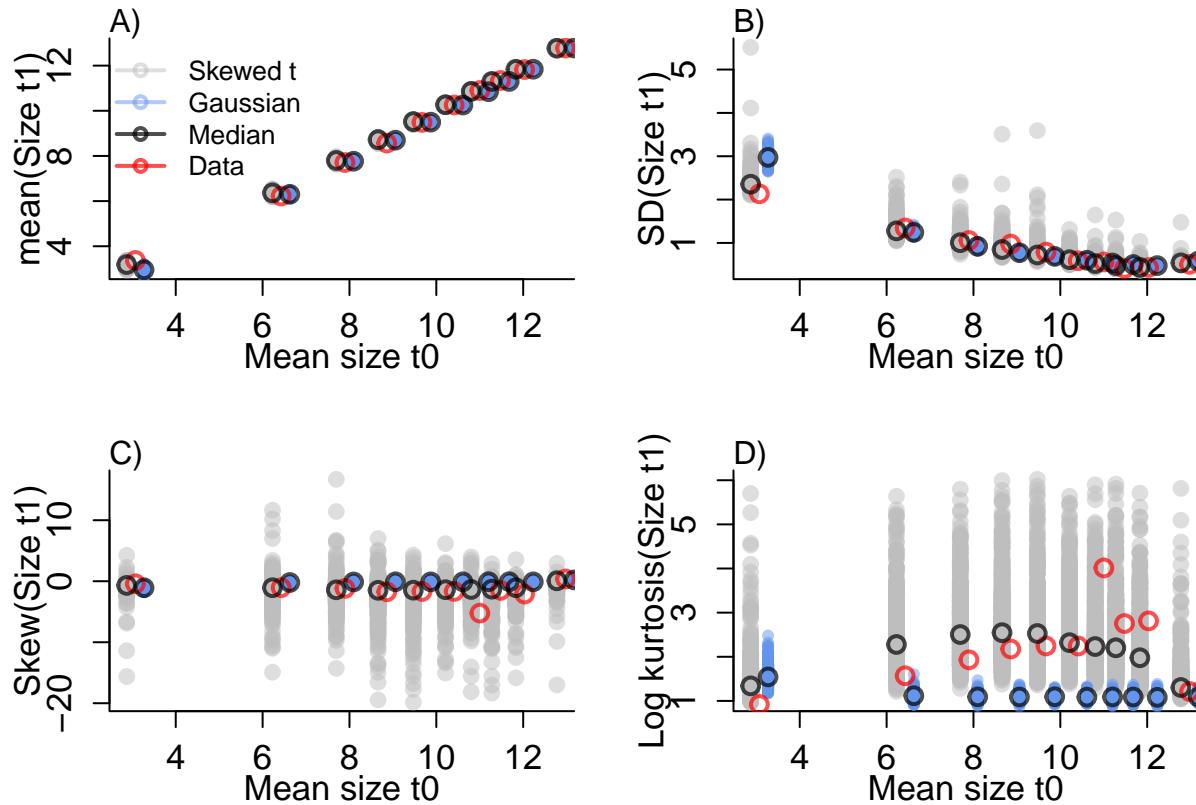
```
CYIM_ST_kernel <- matrix(NA, size_dim, size_dim)
for (i in 1:size_dim) {
  mu_size <- mean(coefs[c(years, plots)]) + coefs["log(vol_t)"] * size_dum[i]
  CYIM_ST_kernel[, i] <- dST3(x = size_dum, mu = mu_size, sigma = exp(coefs["sigma_b0"] +
    coefs["sigma_b1"] * mu_size + coefs["sigma_b2"] * mu_size^2), nu = exp(coefs["nu_b0"] +
    coefs["nu_b1"] * mu_size), tau = exp(coefs["tau_b0"] + coefs["tau_b1"] *
    mu_size + coefs["tau_b2"] * mu_size^2))
}

levelplot(CYIM_ST_kernel, row.values = size_dum, column.values = size_dum, cuts = 30,
  col.regions = rainbow(30), xlab = "log Size t", ylab = "log Size t+1", main = "Skewed t",
  panel = function(...) {
  panel.levelplot(...)
  grid.points(log(CYIM$vol_t), log(CYIM$vol_t1), pch = ".", gp = gpar(cex = 3,
    col = alpha("black", 0.5)))
})
})
```



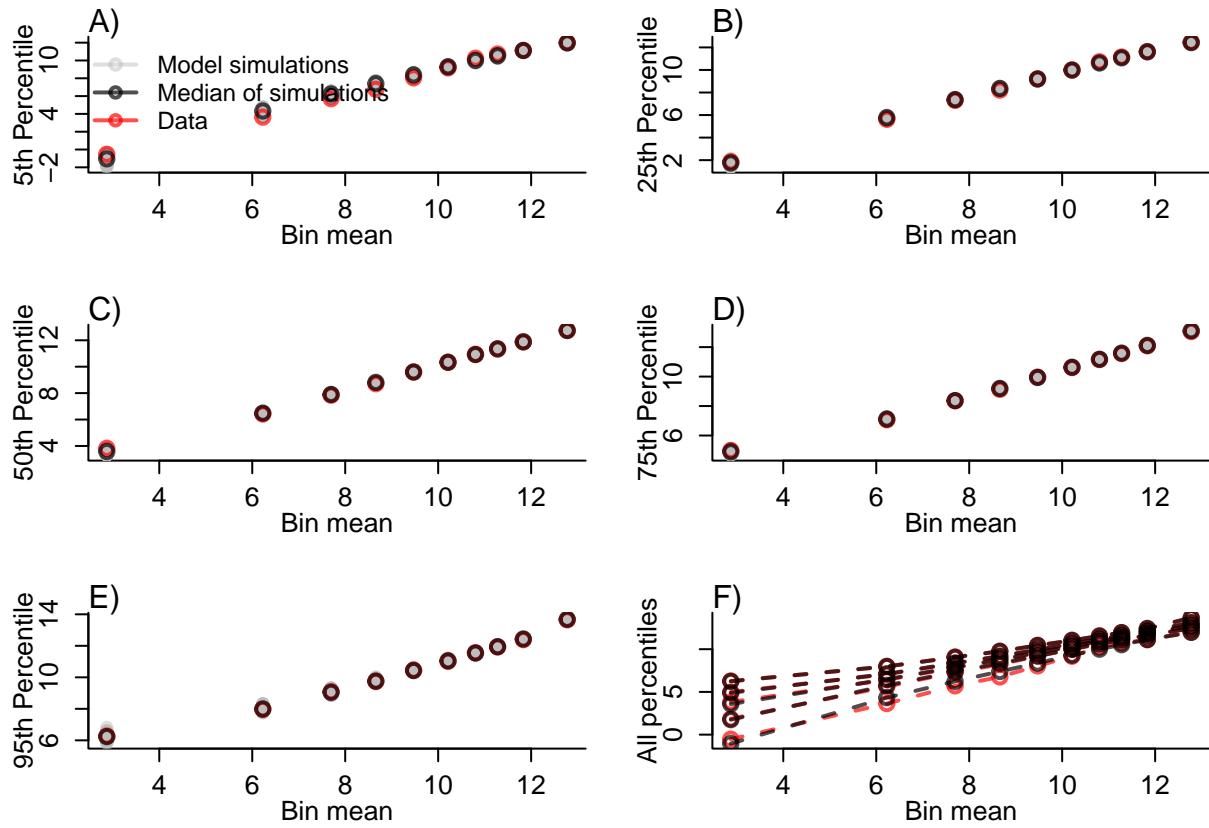
Now finer diagnostics comparing moments and quantiles of the real data against data generated by the fitted models. I am including the Gaussian and ST models for comparison. Now that I know a quadratic term for sigma improves fit, it might be fairer to included this in the Gaussian model, but proceeding with the original form of non-constant variance for now.

The skewed t model generally looks to be fitting well. Is it a strikingly better fit than the Gaussian? In mean and SD, no, which is not surprising because non-constant variance in lme4 can generally handle those. The fitted skewness of the ST is not particularly impressive, at least in this visualization, because the model can generate data with a really wide range of skewness just by chance, including positive and negative skew. Not sure yet if this is a problem, or just a natural consequence of the skewed t or skewed distributions generally. The median skewness estimates look good, i.e., show negative skew consistent with the data. Kurtosis looks great and this is a clear ‘winner’ over the Gaussian.



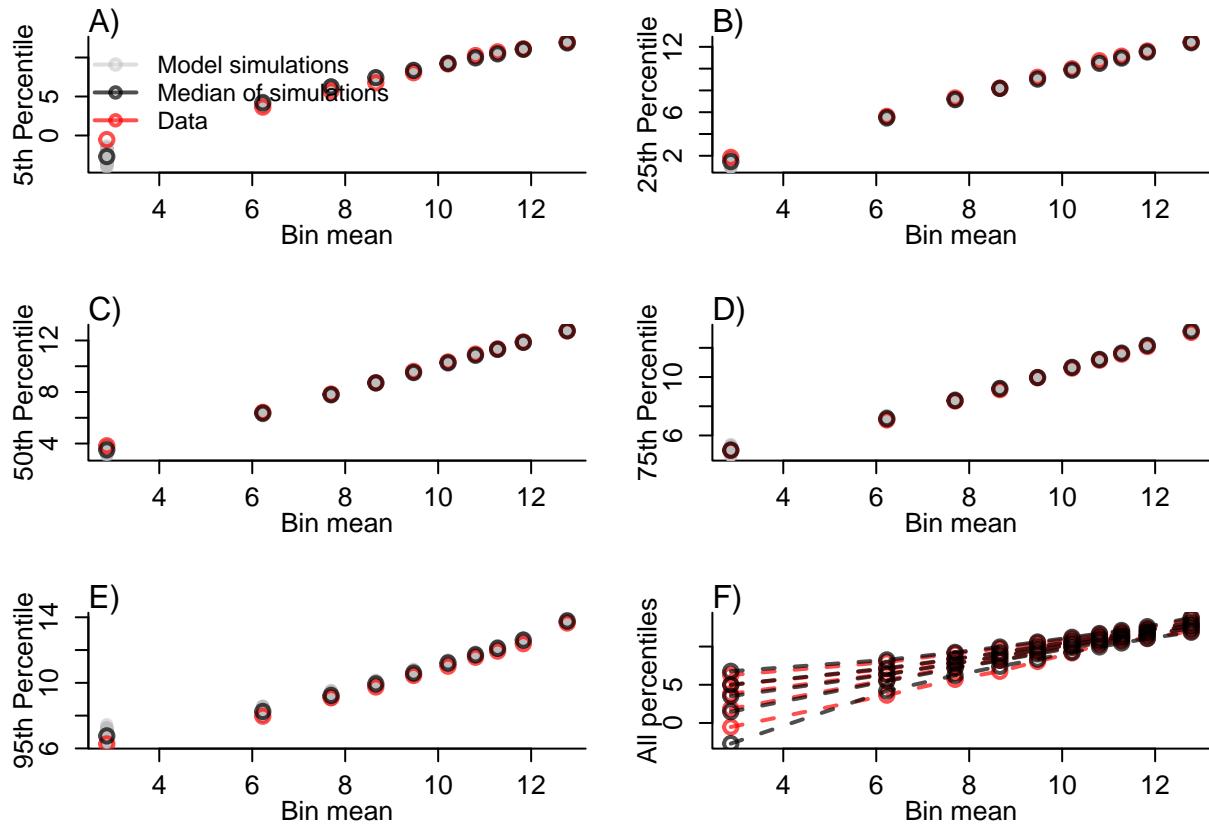
Steve wrote this nifty function to do something similar but with quantiles. I have not added the Gaussian comparison but I can print this out separately. Here is the `skewet t`.

```
quantileComparePlot(log(CYIM$vol_t), log(CYIM$vol_t1), cactus_sim, n_bins)
```



And here is the Gaussian. These look *really* similar, except perhaps a mismatch at the smallest sizes.

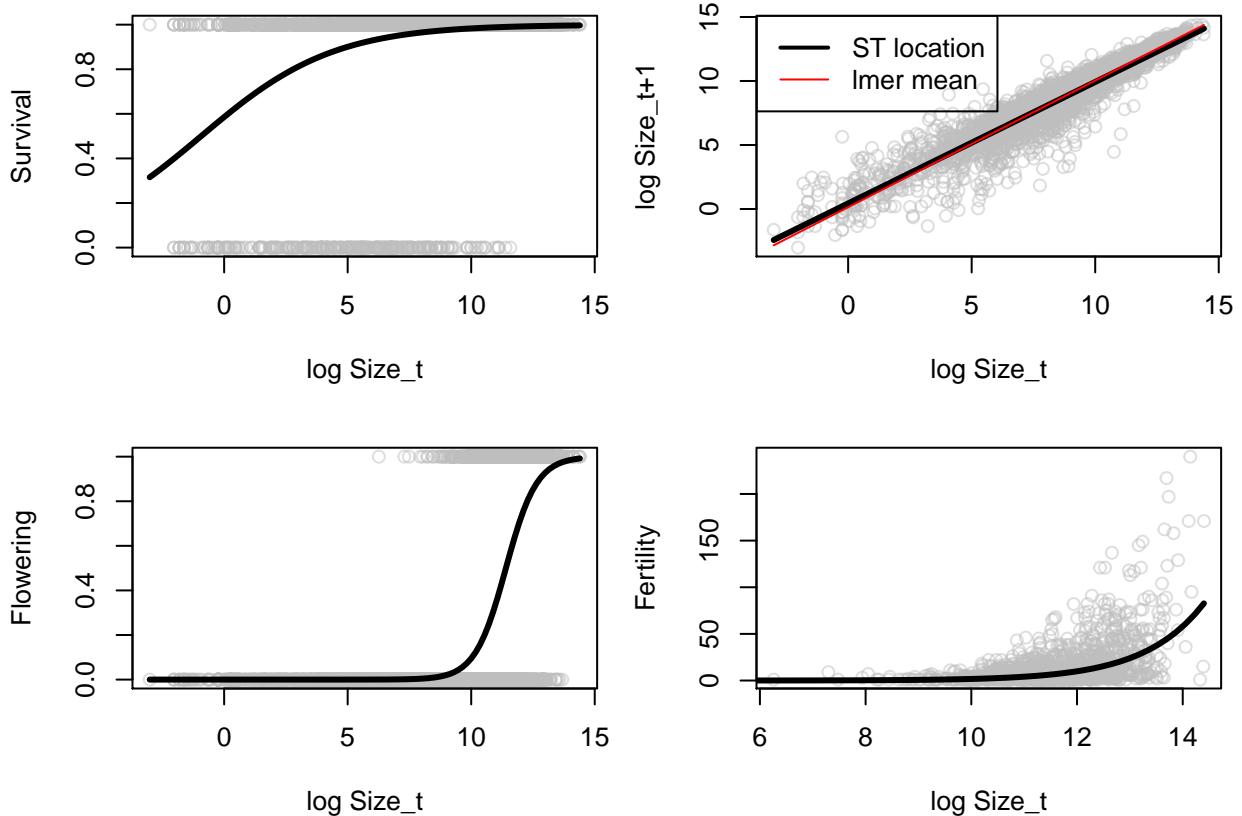
```
quantileComparePlot(log(CYIM$vol_t), log(CYIM$vol_t1), cactus_sim_norm, n_bins)
```



IPM comparisons

Finally, we can compare IPM predictions between the best Gaussian model and the skewed t model.

Here are the size-dependent functions for survival, flowering, and flowerbud production, fit in lme4 with year and plot random effects.



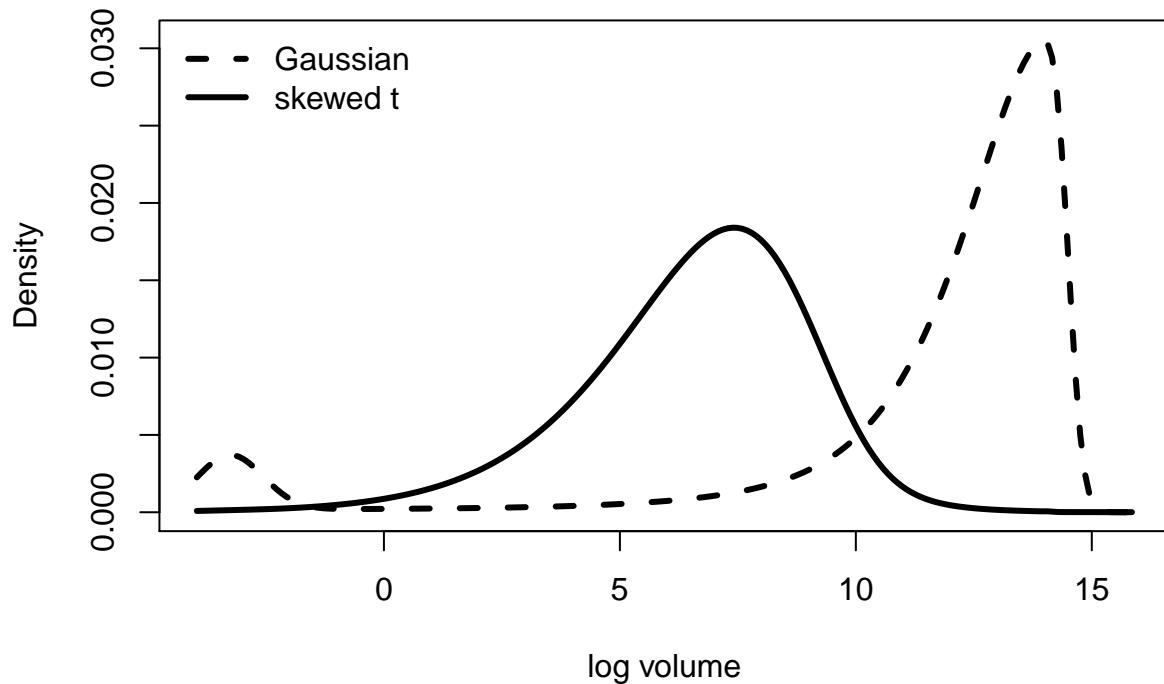
And finally for IPM results. The mean model (averaging across years and plots) predicts very different growth rates:

```
tibble(Growth_Dist = c("Skewed t", "Gaussian"), lambda = c(round(lambda(kernel_ST$IPMmat), 4), round(lambda(kernel_norm$IPMmat), 4)))

## # A tibble: 2 x 2
##   Growth_Dist     lambda
##   <chr>          <dbl>
## 1 Skewed t      0.919
## 2 Gaussian      0.993
```

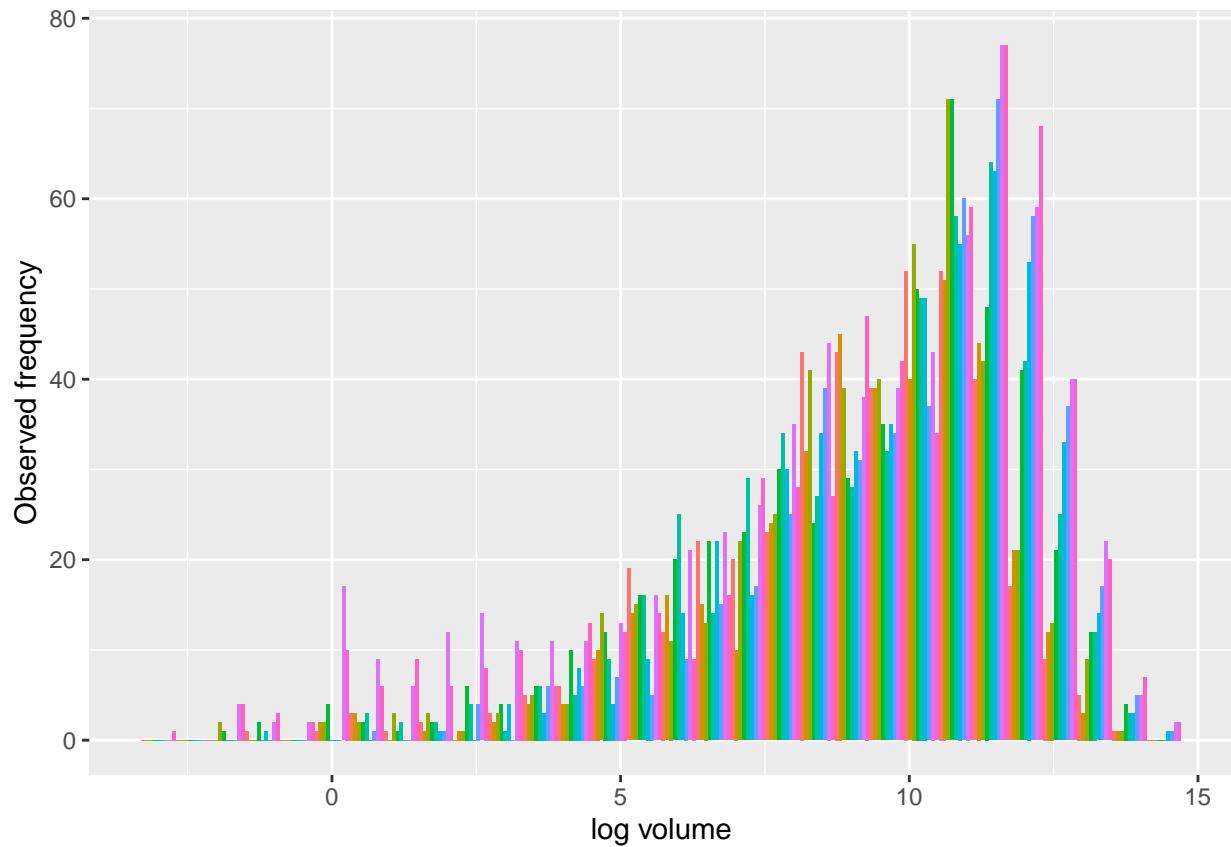
Also, the SSD's are very different and the observed distribution is maybe in between the two of them (in the data plot, colored bars are years). The Gaussian ssd is bi-modal but the lower mode disappears with the skewed t growth kernel, probably because the big reproductive plants are rare in the ssd so you lose the signal of recruitment.

```
ssd_ST <- stable.stage(kernel_ST$IPMmat)[3:(mat.size + 2)]/sum(stable.stage(kernel_ST$IPMmat)[3:(mat.size + 2)])
ssd_norm <- stable.stage(kernel_norm$IPMmat)[3:(mat.size + 2)]/sum(stable.stage(kernel_norm$IPMmat)[3:(mat.size + 2)])
plot(kernel_ST$meshpts, ssd_norm, type = "l", lty = 2, lwd = 3, xlab = "log volume",
     ylab = "Density")
lines(kernel_norm$meshpts, ssd_ST, type = "l", lwd = 3)
legend("topleft", c("Gaussian", "skewed t"), lty = c(2, 1), lwd = 3, bty = "n")
```



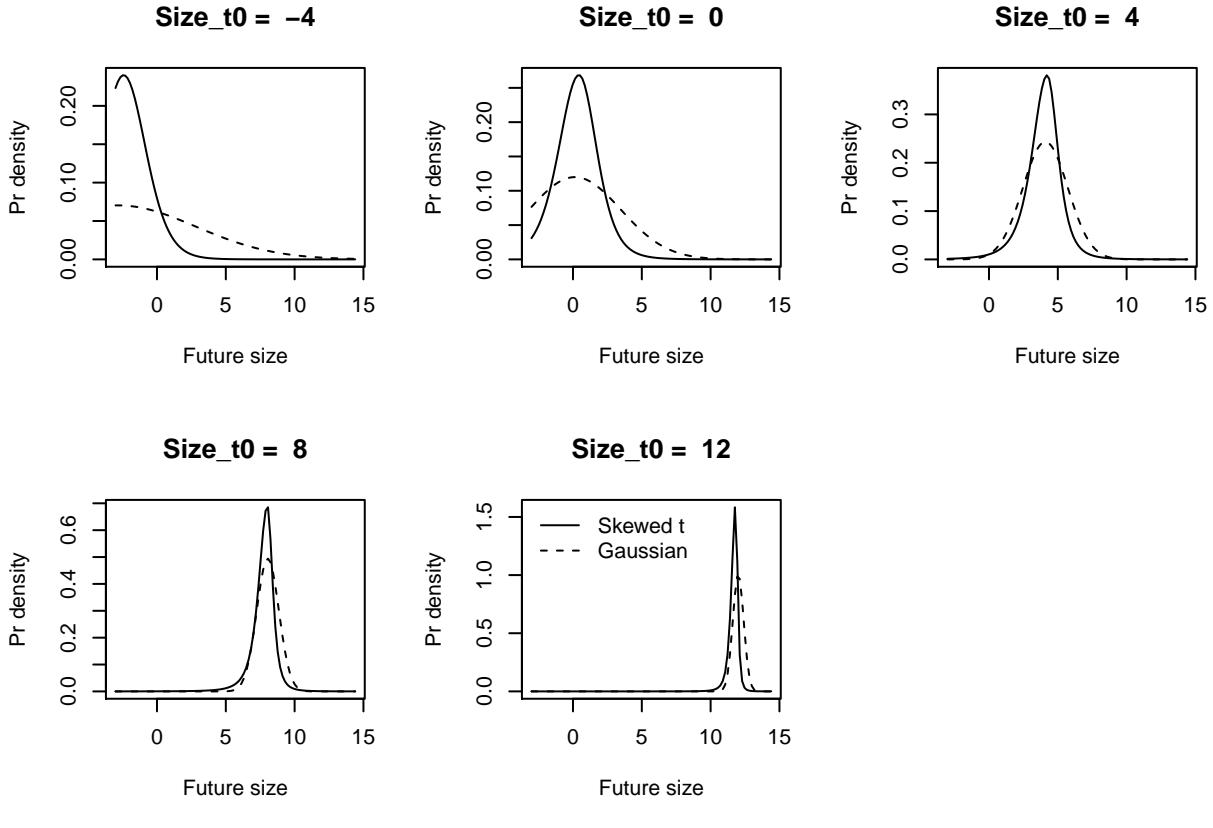
```
CYIM %>% ggplot() + geom_histogram(aes(x = log(vol_t), fill = as.factor(year_t)),
position = "dodge") + theme(legend.position = "none") + xlab("log volume") +
ylab("Observed frequency")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



The difference in SSD's is pretty striking, so I just wanted to have a closer look at the two growth kernels. The skewed t is "peak-ier" than the Gaussian and this causes the Gaussian to allow for large increases in size that don't happen with the skewed t . I think that is where the big difference in SSD comes from.

```
x = c(-4, 0, 4, 8, 12)
par(mfrow = c(2, 3))
for (i in 1:length(x)) {
  plot(size_dum, gxy_ST(x = x[i], y = size_dum, params = cactus_params), type = "l",
       main = paste("Size_t0 = ", x[i]), xlab = "Future size", ylab = "Pr density")
  lines(size_dum, gxy_norm(x = x[i], y = size_dum, params = cactus_params), lty = 2)
}
legend("topleft", legend = c("Skewed t", "Gaussian"), lty = c(1, 2), bty = "n")
```



Conclusions

- The cactus size transition data are not Gaussian, even though a Gaussian might pass the initial smell test, especially with non-constant variance
- The skewed t provides a better fit to the data, most noticeably in terms of kurtosis. There is a negative skew in the data that is reflected in the median estimates of the skewed t simulations, but these simulations produced surprisingly wide variation in skewness. Choice of the skewed t was a little arbitrary because other distributions were also in contention. This could (should?) have been an AIC model selection exercise.
- Fitting models with multiple random effects via shrinkage is tedious. The shrinkage variance estimates for plot is not exactly right because I could not figure out the SE of the reference plot estimate. There is probably a better way to do this. Still, I think showing an example with multiple RFX is worthwhile because this is a common thing that people deal with.
- IPM predictions are very different between the skewed t and Gaussian growth kernels, with all else equal. The tails of the Gaussian allow for large increases in size that the skewed t says are highly unlikely. As a result, they have very different SSD's, with the skewed t SSD much smaller and therefore declining at a faster rate.
- The observed size distribution is not consistent with either SSD. However, all the evidence from this system suggests that the population is declining from previously high abundance (possibly driven by land use change), so it is plausible that it is not at SSD.